

Distributed Semantic Search [Task Outline]

Repo: <https://github.com/akhildhiman7/Distributed-Semantic-Search>

Project Idea: istributed Semantic Search Engine [Project Idea]

| | A | B | C | D | E |
|----|----------|-------------|---------------------------------|---------|----------------------------------|
| 1 | Member | Member Name | Task | Status | Comments |
| 2 | Member 1 | Akhil | E.g. Task 1 - short description | Blocked | E.g, I'm blocked on Akhil task 2 |
| 3 | | | E.g. Task 2 | To Do | |
| 4 | Member 2 | | | To Do | |
| 5 | | | | To Do | |
| 6 | Member 3 | | | To Do | |
| 7 | | | | To Do | |
| 8 | Member 4 | | | To Do | |
| 9 | | | | To Do | |
| 10 | Member 5 | | | To Do | |
| 11 | | | | To Do | |
| 12 | | | | To Do | |

Stack (best-in-class, fast to implement):

- **Storage/Index:** Milvus 2.4+ (HNSW / IVF_FLAT), **MinIO** (object store), **etcd** (metadata)
- **Orchestration:** Docker Compose (2–3 Milvus query nodes + proxy) — faster than K8s, still realistic
- **Embeddings:** Sentence-Transformers all-MiniLM-L6-v2 (384-d; CPU is fine), batched inference
- **API:** FastAPI + pymilvus
- **Dataset:** Kaggle arXiv metadata (title + abstract), target **≥1 GB** text
- **Observability:** Prometheus + Grafana (containerized)
- **Benchmarks:** Locust (QPS/latency), Python harness for cold/hot latency + recovery tests
- **Repo structure:**

```
/infra      # compose files, MinIO/etc, Milvus, Prom/Grafana
/data       # scripts to fetch/clean arXiv, schema
/embed      # embedding pipeline & exporters
/api        # FastAPI service (search/insert/health)
/bench      # load tests, reports
/docs       # VLDB paper, slides, architecture diagram
```

System Architecture (high level)

```
[User/Client]
  -> FastAPI (/search, /insert, /health)
  -> Milvus Proxy (LB)
```

```
-> Milvus Query Nodes (HNSW/IVF index shards; replicas)
-> MinIO (vectors/segments) + etcd (cluster metadata)
[Prometheus] -> scrapes FastAPI & Milvus metrics -> [Grafana Dashboards]
```

Data model (Milvus collection):

- paper_id INT64 (PK, auto_id=false)
- vector FLOAT_VECTOR(384)
- title VARCHAR(512)
- abstract VARCHAR(4096)
- categories VARCHAR(256)
- **Index:** HNSW {M: 16, efConstruction: 200} (fast queries); alt: IVF_FLAT {nlist: 4096} for large batches
- **Metric:** cosine (IP) or L2 (pick 1 and keep consistent)

API Contract (FastAPI):

- POST /search { "query": string, "top_k": int=5, "filters": {"categories": ["cs.LG"]} }
- POST /insert { "paper_id": int, "title": str, "abstract": str, "categories": str } (server embeds & upserts)
- GET /health → { "api": "ok", "milvus": "ok", "index_loaded": true }
- GET /metrics (Prometheus exposition)

Member-wise Detailed Plan

👤 Member 1 — Data Engineer (arXiv Fetch & Clean)

Objective: Deliver clean, deduplicated, ready-to-embed dataset ≥1 GB with robust provenance.

Concrete tasks:

1. Ingestion

- Pull arxiv-metadata-oai-snapshot.json locally.
- Stream-parse JSONL; extract id, title, abstract, categories. Drop empty/short abstracts.

1. Cleaning

- Normalize whitespace; strip HTML/LaTeX; collapse multiple spaces.
- Concatenate title + ". " + abstract → text.
- Deduplicate by normalized title hash + first 200 chars of abstract.

1. Partitioning

- Write **Parquet** partitions (/data/out/clean/part-*.parquet) ~100–250 MB each.
- Produce a **10k row sample** for early integration (sample.parquet).

1. Data dictionary & stats

- CSV of category counts; descriptive stats (avg length, #records).
- Document exact filters so the dataset is reproducible.

Deliverables / DoD

- `data/clean_arxiv_parquet/` (≥ 1 GB total), `data/sample.parquet` (10k)
- `data/README.md` (commands, schema, filters)
- `data/profile.json` (counts, lengths, categories)

Can start Day 1; no dependencies.

👤 Member 2 — ML Engineer (Embeddings Pipeline)

Objective: Produce high-quality embeddings with batched CPU inference; export aligned with metadata.

Concrete tasks:

1. Model & batching

- Use `sentence-transformers all-MiniLM-L6-v2`, `batch_size=64` (tune by RAM).
- Persist embeddings in **NumPy memmap** to avoid RAM blowups.

1. Processing pipeline

- Read partitions sequentially; keep `(paper_id, vector, title, abstract, categories)`.
- Save per-partition **Feather/Parquet + .npy or .npy.memmap**.

1. Integrity & speed

- Hash checks to ensure row order alignment.
- Log `embeddings/speed_report.md` (docs/sec, ETA).

1. Optional optimizations

- Try `normalize_embeddings=True` (cosine).
- Evaluate int8 quantization (optional note in report).

Deliverables / DoD

- `embed/embeddings/part-*.npy` (+ matching metadata parquet)
- `embed/embedding_pipeline.py` (idempotent)
- `embed/README.md` (hardware used, speed, parameters)

Depends on M1's sample early (Day 2), full set by Day 3–4.

👤 Member 3 — Systems Engineer (Milvus + Storage + Indexing)

Objective: Stand up a **multi-node Milvus** over Docker Compose with MinIO & etcd; load data; build index; prove HA.

Concrete tasks:

1. Infra bring-up

- o Compose stack: etcd, MinIO, Milvus standalone → **then** switch to **cluster with 1 proxy + 2 query nodes**.
- o Persist volumes; .env for ports/credentials.

1. Collection & index

- o Create `papers` collection w/ schema above.
- o Insert metadata + vectors in batches from M2 outputs.
- o Build **HNSW** index; `collection.load()`.

1. HA/replication

- o Add second query node; verify proxy LB.
- o Kill one query node mid-search → demonstrate continued service.

1. Exportable scripts

- o `infra/scripts/create_collection.py`, `load_data.py`, `build_index.py`, `smoke_search.py`.

Deliverables / DoD

- `infra/docker-compose.yaml` (cluster mode), `infra/.env.example`
- `infra/scripts/*` (create, load, index)
- `infra/ops_guide.md` (start/stop, failure demo)
- Evidence: screenshot/logs of node loss + continued query success

Can start Day 1 using synthetic vectors; integrate real embeddings by Day 4–5.



Member 4 — Backend Engineer (FastAPI + Client)

Objective: Ship a clean, documented API for `/search` & `/insert` with unit tests and a minimal UI.

Concrete tasks:

1. App skeleton

- o FastAPI app; config via env (`MILVUS_HOST`, `COLLECTION_NAME`, metric).
- o Dependency for SentenceTransformer (for query embedding only).

1. Endpoints

- o POST `/search` → embed query → Milvus search (`top_k`, optional `categories` filter). Return list of `{paper_id, title, abstract, categories, score}`.
- o POST `/insert` → embed & upsert one paper (calls Milvus insert).
- o GET `/health` + `/metrics` (Prometheus client).

1. Quality

- o Unit tests with `pytest` + Milvus mocked interface (adapter class).
- o Rate limit middleware (basic).
- o CORS for simple web client.

1. Demo UI (nice to have)

- Streamlit page in `/api/ui/` with a search box and result list.

Deliverables / DoD

- `api/main.py`, `api/requirements.txt`, `api/Dockerfile`
- `api/tests/` (search & insert tests)
- `api/README.md` (run locally; curl examples)
- Successful integration against M3's Milvus proxy

Can start Day 1 with mock; switch to live Milvus Day 4.



Member 5 — DevOps & Evaluator (Monitoring + Benchmarks + Paper)

Objective: Provide hard numbers, dashboards, and final paper/slides; prove scale and resilience.

Concrete tasks:

1. Observability

- Prometheus scrape for FastAPI (`/metrics`) & Milvus (use exporter or scrape proxy stats).
- Grafana dashboards: latency (P50/P90/P95), QPS, CPU/mem, node up/down.

1. Benchmarks

- **Locust** workload: configurable RPS, query sets of 100 canned queries.
- Python harness to measure cold vs. warm latency, `top_k` sensitivity, index type (HNSW vs IVF) comparison.
- **Failure drill:** kill one query node under load, chart error rate & recovery time.

1. Documentation

- `bench/report.md` with charts (export PNGs from Grafana).
- Assemble **VLDB 4-page** report (Intro, System, Implementation, Evaluation, Discussion, Refs).
- Build crisp 6–8 slide deck; include architecture & dashboards.

Deliverables / DoD

- `infra/prometheus.yml`, `infra/grafana/` provisioning
- `bench/locustfile.py`, `bench/harness.py`
- `docs/COMP6231_Final_Report.(tex/pdf)`, `docs/slides.pdf`
- Dashboard JSON exports + screenshots

Can start Day 1 (wire to mock endpoints), go live once M4 connects to M3.

Integration & Quality Gates

Gate A (Day 4):

- M1 sample + M2 sample embeddings ready
- M3 cluster alive w/ test data; index created
- M4 /search returns real results on sample
- M5 has first Grafana metrics showing API latency

Gate B (Day 9):

- ≥ 1 GB text processed & indexed
- /search stable under Locust at target QPS
- Dashboard shows P95 latency & node health

Gate C (Day 12):

- Failure drill (kill 1 query node): <10s recovery, no API crash
- Side-by-side HNSW vs IVF comparison chart

Final (Day 14):

- Demo script rehearsed; report & slides finalized
-

Acceptance Criteria (Definition of Done)

- **Functionality:** /search returns top-k semantically relevant arXiv papers in <120ms P95 on sample queries; <300ms P95 on full set (targets adaptable by hardware).
 - **Scale:** At least 1 GB text ingested; **index built and loaded**.
 - **Resilience:** Under steady RPS, a node failure **does not** crash API; recovery demonstrated & measured.
 - **Observability:** Grafana dashboards show latency, QPS, node status.
 - **Reproducibility:** Fresh clone + docker compose up + one command to index sample should yield a working demo.
 - **Docs:** VLDB-style 4-pager + slides with metrics & architecture.
-

Two-Week Gantt-Style Schedule (Who/What/When)

| | Day | M1 – Data | M2 – Embeddings | M3 – Milvus/Infra | M4 – API/UI | M5 – DevOps/Bench |
|----|-----|--|--|--|--|--|
| 1 | 1 | Fetch JSONL; schema plan | Set up ST model; test batch | Compose (etcd/MinIO/Milvus-standalone); smoke | FastAPI skeleton; mock Milvus | Prom+Grafana stack; seed dashboards |
| 2 | 2 | Clean, dedup, partition; <code>sample.parquet</code> | Dry-run on sample; memmap setup | Switch to cluster (proxy + 2 query nodes) | Wire /search to mock; unit tests | Connect Prom → API; synthetic load |
| 3 | 3 | Deliver ≥1GB parquet set | Batch encode sample; export vectors+meta | Create collection; load sample vectors | Connect to live Milvus; first real results | First latency charts; basic Locust run |
| 4 | 4 | Data README + profile | Start full embedding run | Build HNSW index; <code>collection.load()</code> | /insert endpoint; CORS | Dashboard polish; capture baseline |
| 5 | 5 | (buffer) | Full run continues | Bulk insert automation | Response shaping; error handling | Draft benchmark plan (QPS tiers) |
| 6 | 6 | (buffer) | Finish full encode; export | Load full vectors; rebuild index | End-to-end validation on full set | Baseline run on full set (charts) |
| 7 | 7 | (buffer) | Speed report + params | Tune ef, M (HNSW); alt IVF build | Small UI (Streamlit) | Report skeleton; import figures |
| 8 | 8 | (support) | (support) | Add 3rd query node; verify LB | Logging, /health, /metrics | Load test (RPS ramp, P95/P99) |
| 9 | 9 | (support) | (support) | Optimize insert/search params | API hardening; timeouts | Compare HNSW vs IVF; table/plots |
| 10 | 10 | (support) | (support) | Failure drill: kill node under load | Graceful errors; retries | Recovery chart; error budget notes |
| 11 | 11 | (support) | (support) | Stabilize configs; ops guide | Final API README; curl scripts | Draft findings; finalize dashboards |
| 12 | 12 | (support) | (support) | Freeze infra; tag release | Freeze API; tag release | Write VLDB report (Methods/Results) |
| 13 | 13 | (support) | (support) | Demo rehearsal checklist | Demo script; sample queries list | Slides (arch, metrics, drill) |
| 14 | 14 | (all) | (all) | Final demo | Final demo | Final paper + slides |

Notes on “best quality” choices

- **HNSW** chosen for **low latency** at moderate memory; IVF available for ablations.
- **Memmap embeddings** to avoid memory spikes during export/load.
- **Proxy + multiple query nodes** gives real LB & resilience semantics without K8s overhead.
- **Prometheus/Grafana** over ad-hoc prints → professional-grade observability.
- **Contract-first API** + mocks → M4 unblocked from Day 1.

- **Gate-based integration** avoids late surprises; failure drill is non-negotiable.